



# Recommender Systems

---

Machine Learning  
Felipe Alonso-Atienza



Universidad  
Rey Juan Carlos

# Outline

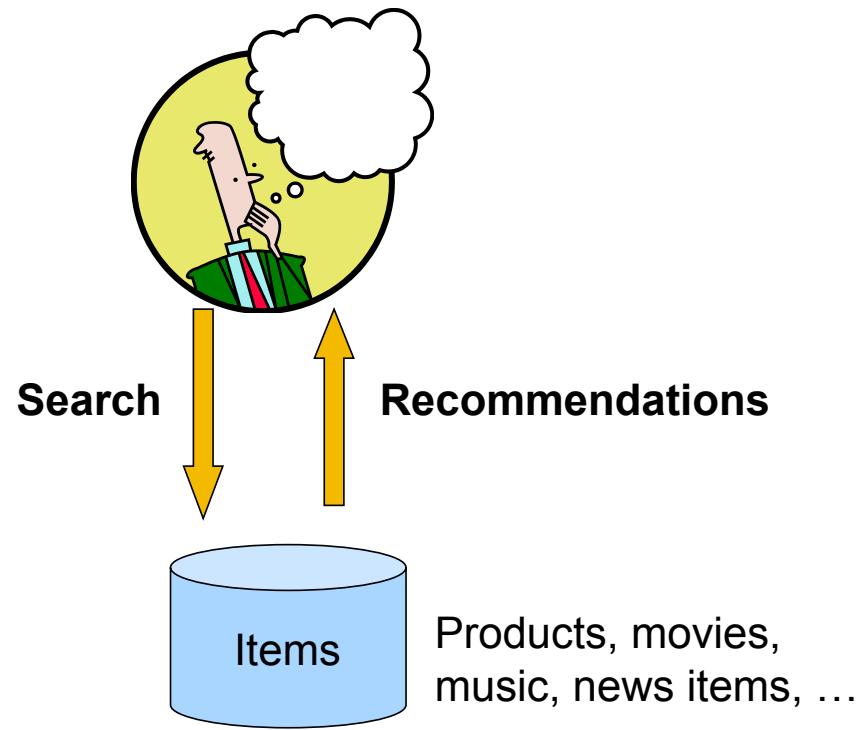
---

1. Introduction
2. Simple recommenders
3. Content-based recommenders
4. Collaborative Filtering
5. Evaluating recommender systems
6. Latent Factors models
7. Recommenders using Deep Learning
8. Recommenders in action (next lecture)



# What's a RS?

---



Given a set of items  $S$  and a set of users  $U$ ,  
we want to recommend item  $i \in S$  to user  $u \in U$

# Introduction

---

- Relatively novel machine learning discipline, which has received low attention in academia during the last years
- However, one of the most popular applications of machine learning in technology companies:
  - Substantial fraction of the revenue
    - Amazon: 35% of total sales
    - [Netflix](#): over 80% of watched TV shows/movies by recommendation
    - [YouTube](#): recommendation algorithm generates more than 70% of views



**EL PAÍS**   **EL MUNDO**



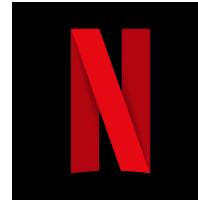
# The origin: the Netflix prize



- From renting DVDs to content creator
- **1M\$ prize**
- Training data
  - 100 million rating,
  - 480.000 users,
  - 17.770 movies
  - 6 years data
- Test data
  - Last few ratings of each user
  - 2.8 million



# The Netflix prize



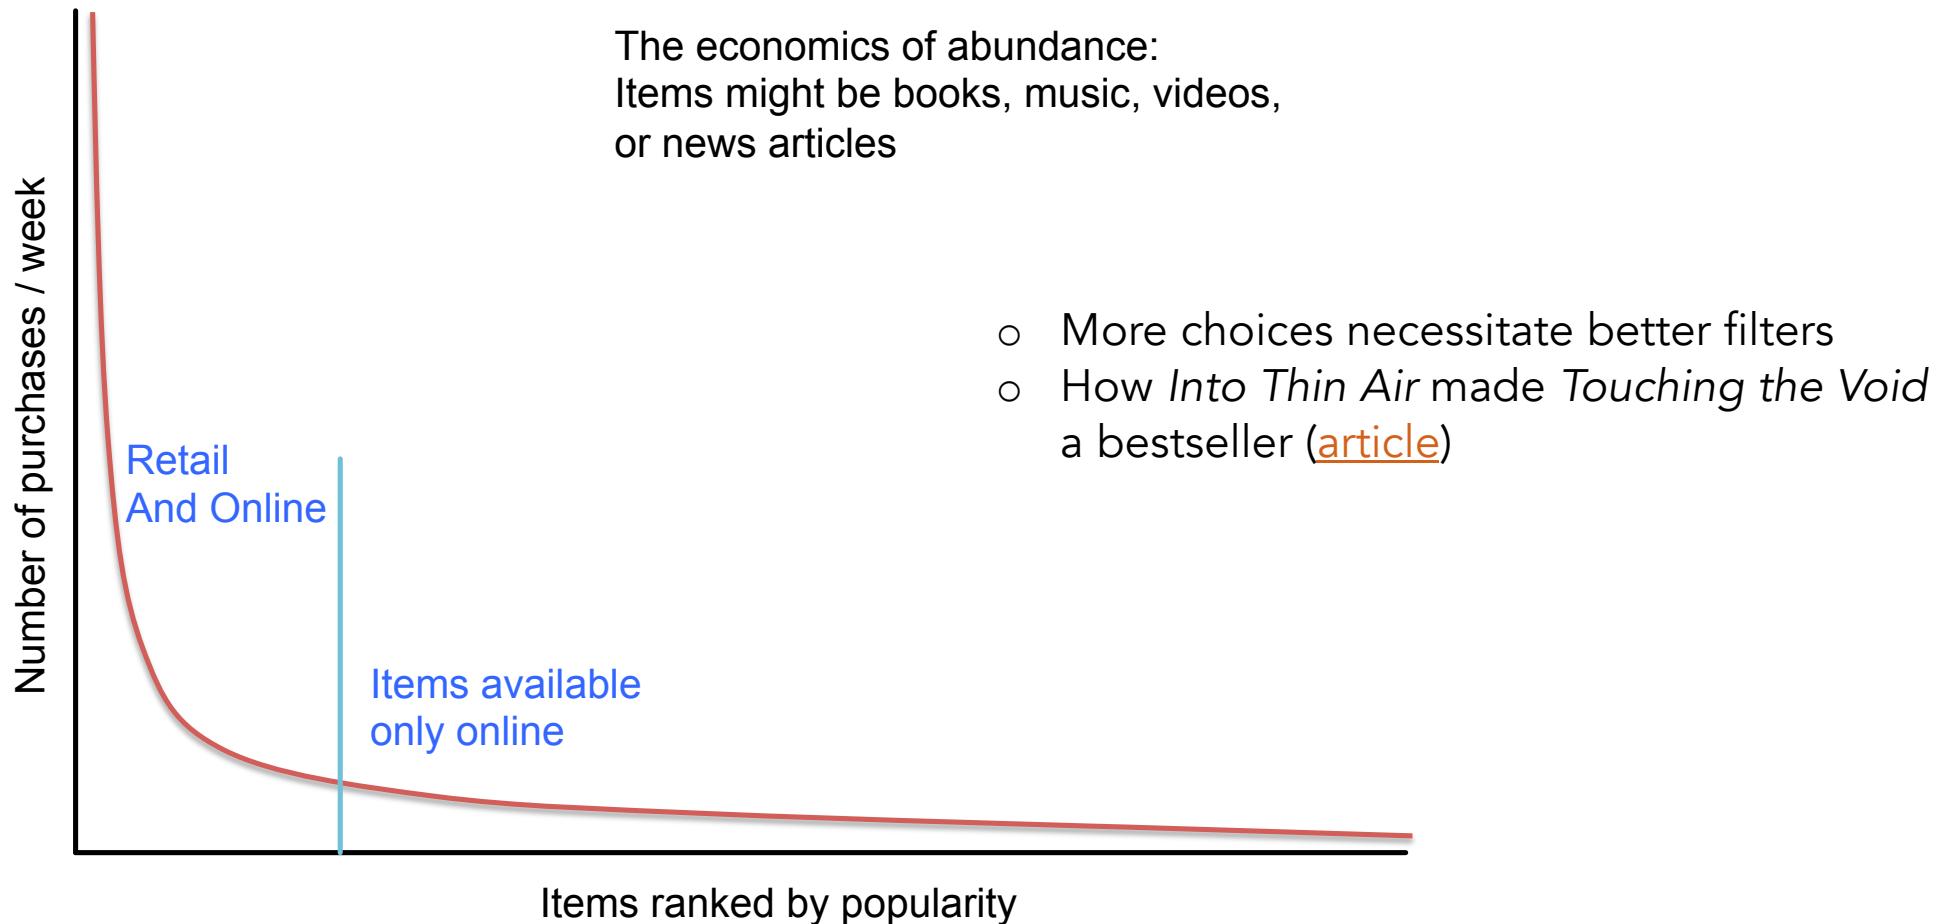
- “The Netflix Prize abstracted the recommendation problem to a simplified proxy of accurately **predicting ratings**. It is now clear that this is just one of many components in an effective industrial recommendation system. They also need to account for factors like **diversity, context, evidence, freshness, and novelty**.”

[Xavier Amatriain et. Al. [Past, Present, and Future of Recommender Systems: An Industry Perspective](#), Recsys '16]



# Motivation: the long tail

---



# Examples

---

- Products/items
  - Global vs session recommendations



- Need of personalized content: coherent and diverse recommendations



- Songs (exploiting music characteristics), radio



- News

**EL PAÍS** **EL MUNDO**

# Examples

---

- People / friends
  - Items are also users



- Search results
  - Pandora vs Pandora

PANDORA



# What's a RS?

---

- Provide a list of (top-N) recommendations instead of predicting ratings
- Amazon does not show your predicted rating

# Outline

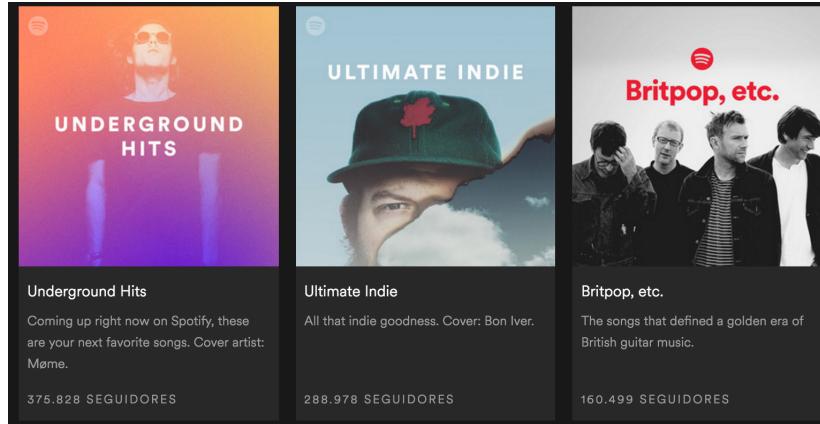
---

1. Introduction
2. Simple recommenders
3. Content-based recommenders
4. Collaborative Filtering
5. Evaluating recommender systems
6. Latent Factors models
7. Recommenders using Deep Learning
8. Recommenders in action (next lecture)

# Non-personalized recommenders

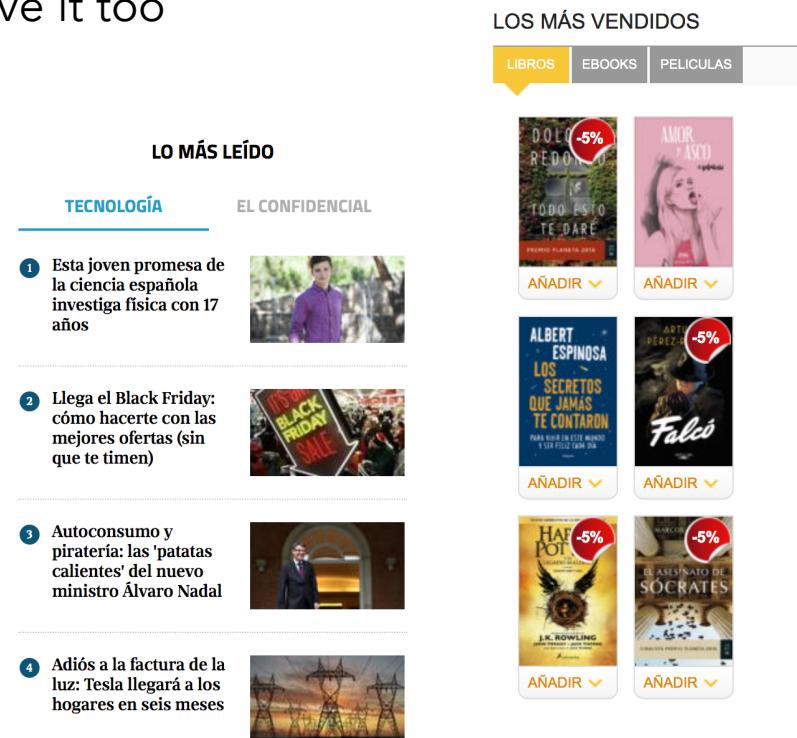
---

- Lists
  - Favorites
  - Essentials



# Non-personalized recommenders

- Recommendations based on popularity
  - Population behavior can be used to guide suggestions
    - “All people love this, so you might love it too”
      - Most read articles...
      - Most watched video...
      - Most played song...
      - Most recent ...
    - Most popular restaurant?
    - Highest rated movie?



# Non-personalized recommenders

---

- Ranking in news aggregators / communities
  - Popularity / age



# Personalized recommendations

---

- How to generate personalized recommendations?
  1. User information/tastes
  2. Item/product characteristics
  3. Interaction between users and items
- How do we know what users would like?
  - Explicit feedback: ask the customer
    - Extra work -> sparse data
    - Different standards
  - Implicit feedback
    - Click data (not reliable, fraud)
    - Purchases (money) / video consumption (time)



# Personalized recommendations

---

- Utility matrix:  $U \times S \rightarrow R$ 
  - $U$  = set of users/customers
  - $S$  = set of Items/products
  - $R$  = set of ratings; e.g., [0-1]; 1-5 stars.
    - What does an empty cell mean?

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4



# Personalized recommendations: Key problems

---

- 1. Gathering ratings:**
  - I. Asking people
  - II. Estimate from user actions (e.g. purchase implies high rating)
- 2. Estimate unknown ratings from the known ones**
  - I. Most people have not rated most items
  - II. Cold start:
    - New items have no ratings
    - New users have no history
- 3. Evaluate the success/performance of recommendation methods**

# Outline

---

1. Introduction
2. Simple recommenders
- 3. Content-based recommenders**
4. Collaborative Filtering
5. Evaluating recommender systems
6. Latent Factors models
7. Recommenders using Deep Learning
8. Recommenders in action (next lecture)

# Content-based systems

- Use properties of the items (metadata) to recommend
  - Movies: genre, year of release, set of actors, ...

Movie	action	adventure	animation	children's	comedy	crime	documentary	drama	fantasy	film-noir	horror	musical	western	mystery	romance	sci-fi	thriller	war	western2
Toy Story	0	1	1	1	1	0	0	0	1	0	0	0	0	0	0	0	0	0	0
Jumanji	0	1	0	1	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0
Grumpier Old Men	0	0	0	0	1	0	0	0	0	0	0	0	0	0	1	0	0	0	0
Waiting to Exhale	0	0	0	0	1	0	0	1	0	0	0	0	0	0	1	0	0	0	0
Father of the Bride	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0

```
def computeYearSimilarity(self, movie1, movie2, years):  
    diff = abs(years[movie1] - years[movie2])  
    sim = math.exp(-diff / 10.0)  
    return sim
```

- $\text{sim\_total} = \text{sim\_genre} * \text{sim\_years}$

<https://sundog-education.com/recsys/>

# Content-based systems

---

- Use properties of the items (metadata) to recommend
  - Movies: genre, year of release, set of actors, ...

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)	$x_1$ (romance)	$x_2$ (action)
Love at last	5	5	0	0	0.9	0
Romance forever	5	?	?	0	1.0	0.01
Cute puppies of love	?	4	0	?	0.99	0
Nonstop car chases	0	0	5	4	0.1	1.0
Swords vs. karate	0	0	5	?	0	0.9

- This item information can be combined with the user profile and the session information:
  - User info: age, purchase history, ...
  - Session information: day of the purchase, time of the purchase, etc...

# Formulation

---

- This problem can be formulated as a classical machine learning problem
- Regression problem
  - Estimate unknown ratings for each user based on item features (and on user and session features)
- Classification problem
  - Determine if a user would like/dislike an item based on item features (and on user and session features)

# Formulation

- Regression problem
  - Estimate unknown ratings for each user based on item features (and on user and session features)

To learn  $\underline{\theta^{(j)}}$  (parameter for user  $j$ ):

$$\rightarrow \min_{\theta^{(j)}} \frac{1}{2} \sum_{i:r(i,j)=1} \left( (\theta^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{k=1}^n (\theta_k^{(j)})^2$$

To learn  $\underline{\theta^{(1)}}, \underline{\theta^{(2)}}, \dots, \underline{\theta^{(n_u)}}$ :

$$\min_{\theta^{(1)}, \dots, \theta^{(n_u)}} \frac{1}{2} \sum_{j=1}^{n_u} \sum_{i:r(i,j)=1} \left( (\theta^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^{(j)})^2$$

$\Theta^{(1)}, \dots, \Theta^{(n_u)}$

From: Andrew Ng. Machine Learning (Coursera)



# Pros & cons

---

- PROS
  1. Personalized recommendations
  2. Features can capture context (time of the day, what I just saw)
  3. Handles missing information (age of the user, limited purchase history)
  4. No need data for other users (items easy to add)
  5. Explanation for recommended items
- CONS
  1. Features might not be available or the are hard to find
  2. Overspecialization:
    - Never recommends items outside user's content profile
    - Unable to exploit quality judgments of other user
  3. Cold-start problem for new users



# Outline

---

1. Introduction
2. Simple recommenders
3. Content-based recommenders
- 4. Collaborative Filtering**
5. Evaluating recommender systems
6. Latent Factors models
7. Recommenders using Deep Learning
8. Recommenders in action (next lecture)

# Collaborative Filtering (CF)

---

- IDEA: exploit information on what other users have done (rated/purchased) on different items

*“People who bought this also bought...”*



- US PATENT 7113917 B2: Personalized recommendations of items represented within a database, 2006

# Collaborative Filtering (CF)

---

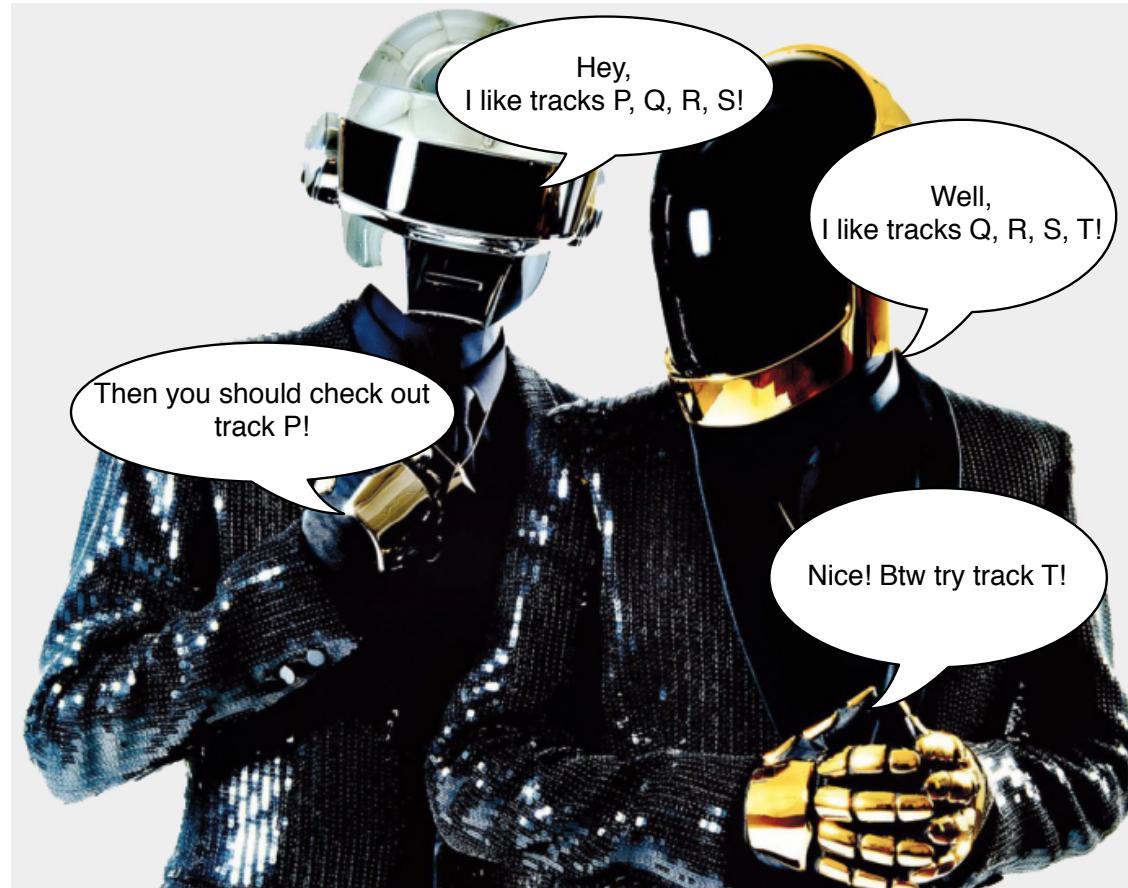


Image via Erik Bernhardsson

# CF taxonomy

---

1. Co-occurrence matrix: "People who bought this also bought..."
  - a. Beers and diapers
2. Memory-based CF
  - a. user-user: "Users who are similar to you also liked ..."
    - If two users have rated similarly a number of movies, then they are probably similar
  - b. item-item: "Users who liked this movie also liked ..."
    - If two movies get similar ratings from lots of users, then they are probably similar
3. Model-based CF, also known as Latent factor models
  - Memory-based CF:
    - Find similarities among users
    - Find similarities among items



# Outline

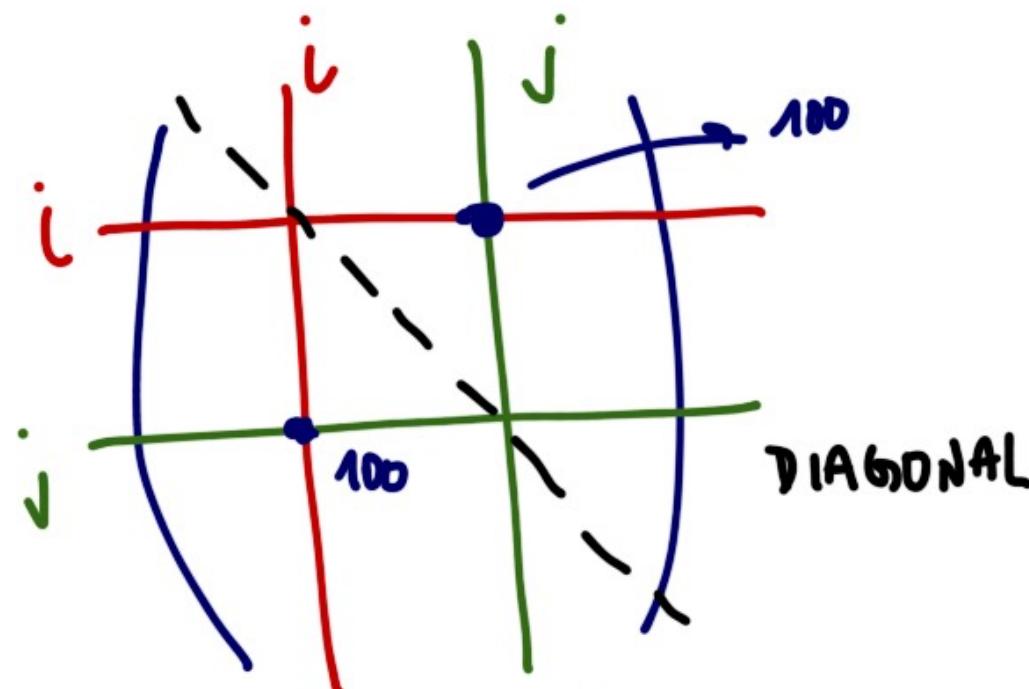
---

1. Introduction
2. Simple recommenders
3. Content-based recommenders
- 4. Collaborative Filtering**
  - I. Co-occurrence matrix
5. Evaluating recommender systems
6. Latent Factors models
7. Recommenders using Deep Learning
8. Recommenders in action (next lecture)



# Co-occurrence matrix

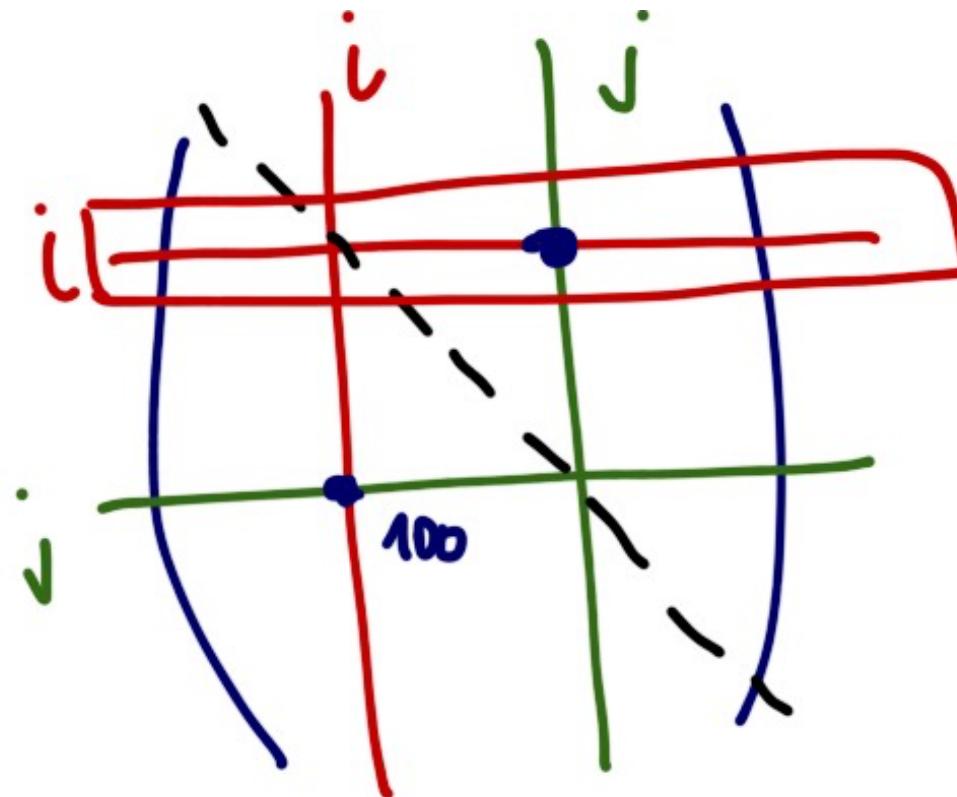
- # of users who purchased both items  $i$  &  $j$ 
  - $i$  = mountain boots
  - $j$  = mountain jacket



# Making recommendations

---

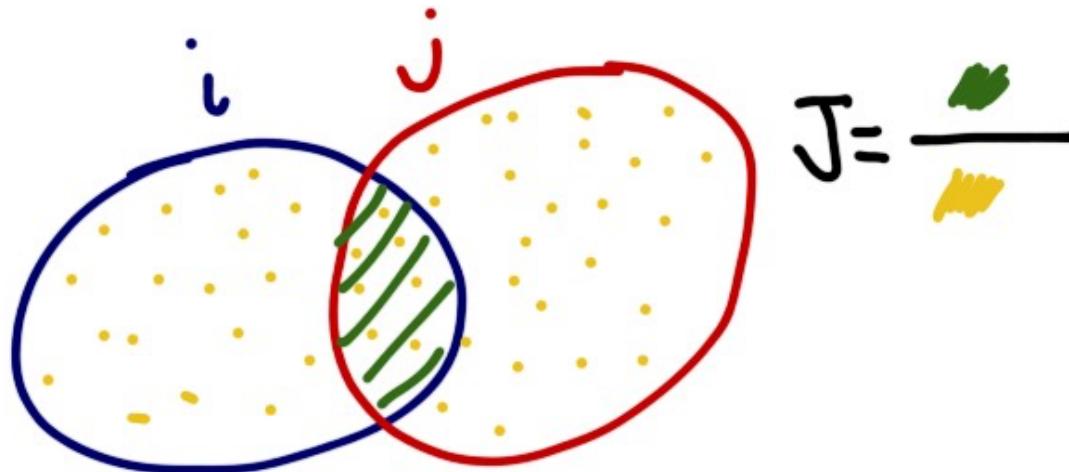
- Look at the specific row  $i$
- Recommend other items with largest counts



# Effect of popular items

---

- What if there are very popular items?
- Normalize co-occurrence Matrix
- Jaccard similarity
  - who purchased **i and j** divided by
  - who purchased **i or j**



# Co-occurrence matrix

---

- Also known as association rules or market basket analysis
  - MLxtend library
- Cons:
  - Only current situation, no history
    - Weighted average of purchases
  - No personalized features

# Outline

---

1. Introduction
2. Simple recommenders
3. Content-based recommenders
4. **Collaborative Filtering**
  - I. Co-occurrence matrix
  - II. **user-user CF**
5. Evaluating recommender systems
6. Latent Factors models
7. Recommenders using Deep Learning
8. Recommenders in action (next lecture)



# user-user CF

---

- “Users who are **similar** to you also liked ...”
- Find similarities among users

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

HP = Harry Potter

TW = Twilight (Crepúsculo)

SW = Star Wars

# user-user CF: Jaccard Similarity

---

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

- $J(A,B) = |r_A \cap r_B| / |r_A \cup r_B| = 1 / 5$
- $J(A,C) = 2 / 4$
- $J(A,B) < J(A,C)$  .... expected?
- KO: ignores rating values!

# user-user CF: cosine distance

---

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

- $\cos(A,B) = r_A \cdot r_B^T / \|r_A\| \cdot \|r_B\| = 5*4 / (\sqrt{42} * \sqrt{66}) = 0.38$
- $\cos(A,C) = 0.32$
- KO: treats missing ratings as negative

# user-user CF: centered cosine

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	2/3			5/3	-7/3		
B	1/3	1/3	-2/3				
C				-5/3	1/3	4/3	
D		0					0

- Normalize ratings by subtracting row mean

# user-user CF: centered cosine

---

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	2/3			5/3	-7/3		
B	1/3	1/3	-2/3				
C				-5/3	1/3	4/3	
D		0					0

- $\cos(A,B) = 0.09$
- $\cos(A,C) = -0.56$
- Treats missing ratings as average
- Also known as Pearson Correlation

# Making recommendations

---

- Let  $r_u$  be the vector of user  $u$ 's ratings
- The prediction for user  $u$  and item  $i$

$$\hat{r}_{ui} = \frac{\sum_{u'} sim(u, u') r_{u'i}}{\sum_{u'} |sim(u, u')|}$$

- Options:
  1. Use only the set of  $k$  users most similar to  $u$
  2. Taking into account user bias

$$\hat{r}_{ui} = \bar{r}_u + \frac{\sum_{u'} sim(u, u') (r_{u'i} - \bar{r}_{u'})}{\sum_{u'} |sim(u, u')|}$$

# Outline

---

1. Introduction
2. Simple recommenders
3. Content-based recommenders
- 4. Collaborative Filtering**
  - I. Co-occurrence matrix
  - II. user-user CF
  - III. item-item CF**
5. Evaluating recommender systems
6. Latent Factors models
7. Recommenders using Deep Learning
8. Recommenders in action (next lecture)



# item-item CF

---

- “Users who liked this movie also liked ...”
- Find similarities among items
- Same procedure as un user-user CF
  - Transpose your utility matrix

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3				3	

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3				3	

# Outline

---

1. Introduction
2. Simple recommenders
3. Content-based recommenders
- 4. Collaborative Filtering**
  - I. Co-occurrence matrix
  - II. user-user CF
  - III. item-item CF
- IV. Global baseline**
5. Evaluating recommender systems
6. Latent Factors models
7. Recommenders using Deep Learning
8. Recommenders in action (next lecture)

# CF + Global baseline

---

- Model local and global effects
- We want to estimate Joe's rating for the movie **The Sixth Sense**
- Global Baseline approach
  - Mean movie rating: **3.7 stars**
  - Joe rates **0.2 stars** below avg.
  - The Sixth Sense is **0.5 stars** above avg.
  - Baseline estimate:  $3.7 + 0.5 - 0.2 = 4$  stars

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{u'} sim(u, u')(r_{u'i} - b_{u'i})}{\sum_{u'} |sim(u, u')|}$$

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_j sim(i, j)(r_{uj} - b_{uj})}{\sum_j |sim(i, j)|}$$

# CF Comparison

---

- **Item-item**
  1. More stable (allows pre-calculations)\*.
  2. Useful when having more users than items.
  3. Easy to justify recommendations
    - “you get a recommendation for movie Y because you liked movie X, which is similar to movie Y”
- **User-user**
  1. Better way to give recommendations
    - Serendipity

\* Linden et al. "[Amazon.com Recommendations Item to Item Collaborative Filtering](#)"

# CF Comparison Pros & cons

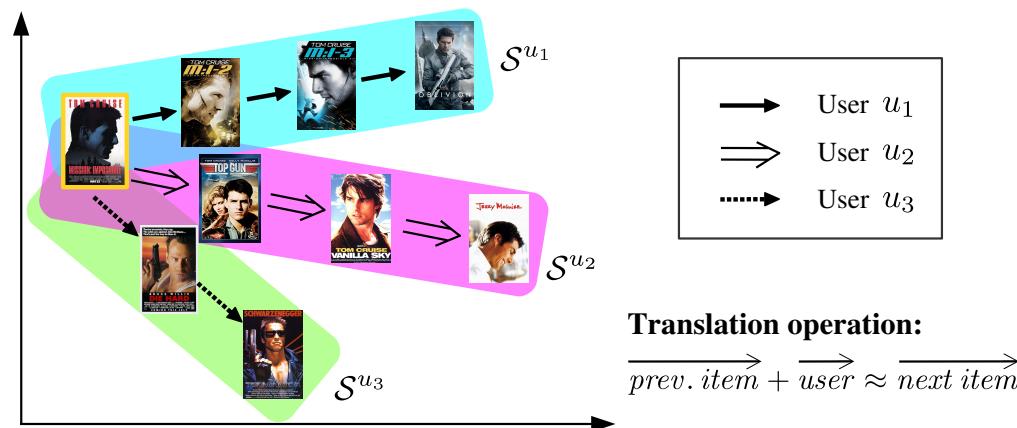
---

- PROS
  1. Easy to compute, specially successful in generating personalized recommendations.
  2. Works for any kind of item
- CONS
  1. Utility matrix is sparse:
    - can lead to bad prediction if the article is unpopular and very few users have given feedback about them
  2. It is difficult to find related users or items for users with extraordinary taste
  3. Cold start (at least 20+ ratings before produce accurate recommendations)



# CF possible variations

- Only consider items with high rating
- Only consider user with high similarity score
- Consider sequence of events:
  - What would be the next movie to watch?



<https://cseweb.ucsd.edu/~jmcauley/pdfs/recsys17.pdf>

# Outline

---

1. Introduction
2. Simple recommenders
3. Content-based recommenders
4. Collaborative Filtering
- 5. Evaluating recommender systems**
6. Latent Factors models
7. Recommenders using Deep Learning
8. Recommenders in action (next lecture)



# Evaluation (GO TO Slide 17)

---

- Off-line
  - Regression: RMSE, MAE
  - Classification:
    - Be on top-N list
    - Rating > threshold
  - Ranking: precision@k, Kendall's tau.
  - Recommenders:
    - Hit rate, average reciprocal hit rate, coverage, diversity
- On-line
  - A/B testing



# Evaluation: RMSE

- Popular metric: RMSE over a test set ( $T$ )

		movies				
		1	3	4		
users	1	3	5			5
	3		4	5		5
			3			
	3					
	3					
	2			2		2
					5	
	2	1				
	3			3		
	1					

		movies				
		1	3	4		
users	1	3	5			5
			4	5		5
	3					
	3					
	2			?		?
			?		?	
	2	1				?
	3			?		
	1					

$$\text{RMSE} = \sqrt{\frac{\sum_{(u,i) \in T} (r_{ui} - \hat{r}_{ui})^2}{N_T}}$$

- KO:
  - lack of prediction context
  - Same penalization for high and low ratings



# Outline

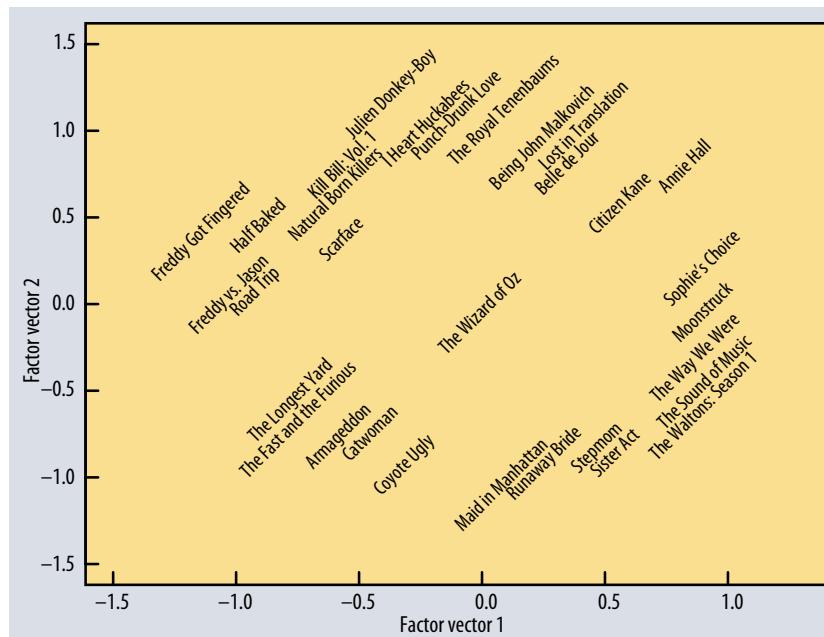
---

1. Introduction
2. Simple recommenders
3. Content-based recommenders
4. Collaborative Filtering
5. Evaluating recommender systems
- 6. Latent Factors models**
7. Recommenders using Deep Learning
8. Recommenders in action (next lecture)

# Latent factors

---

- Idea:
  - Ratings are influenced by a set of factors (e.g. amount of action in movies, comedy vs drama)
  - These factors are not obvious



- The goal is to infer those so called latent factors from the rating data

# Latent factors

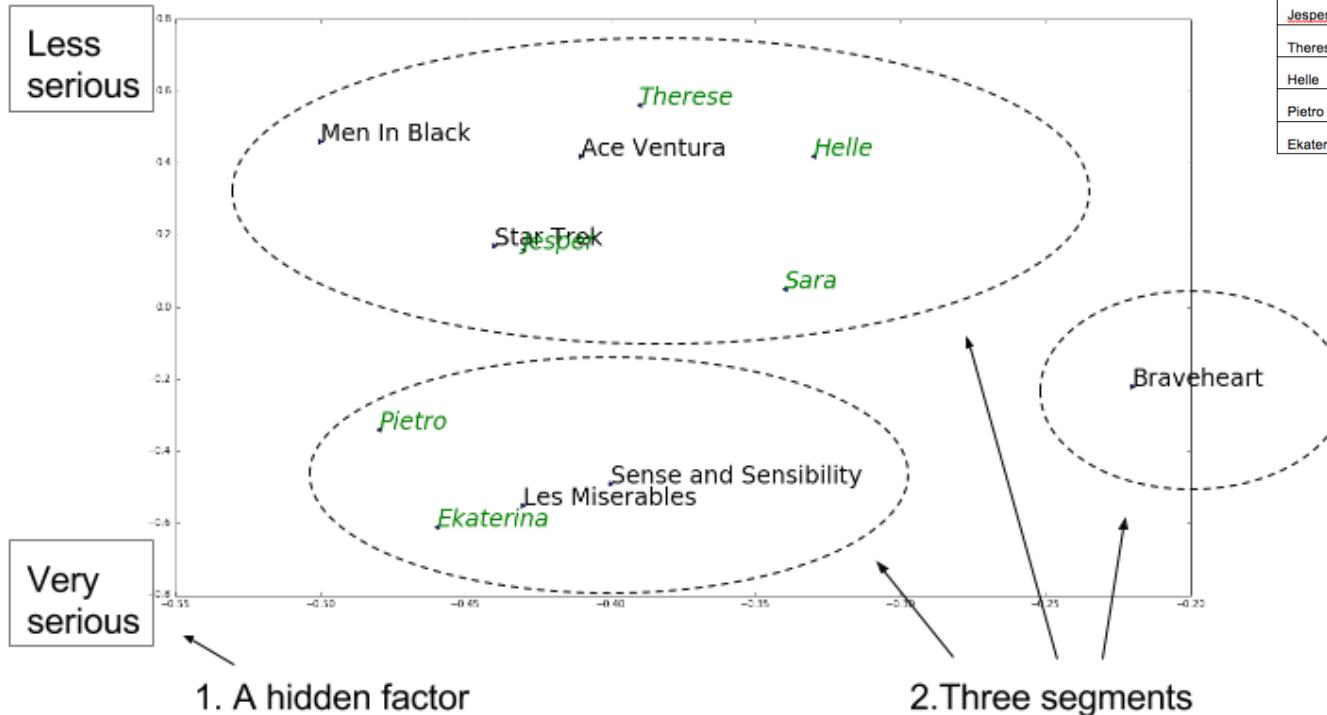
- Idea: matrix factorization



	Comedy	Action	Comedy	Action	Drama	Drama.
Sara	5	3		2	2	2
Jesper	4	3	4		3	3
Therese	5	2	5	2	1	1
Helle	3	5	3		1	1
Pietro	3	3	3	2	4	5
Ekaterina	2	3	2	3	5	5

# Latent factors

- Idea: matrix factorization



	MEN IN BLACK	STAR TREK	ACE VENTURA	BRAVEHEART	SENSE AND SENSIBILITY	LES MISERABLES	LES MISERABLES
	Comedy.	Action	Comedy	Action	Drama	Drama	Drama
Sara	5	3		2	2	2	2
Jesper	4	3	4		3	3	3
Therese	5	2	5	2	1	1	1
Helle	3	5	3		1	1	1
Pietro	3	3	3	2	4	5	5
Ekaterina	2	3	2	3	5	5	5

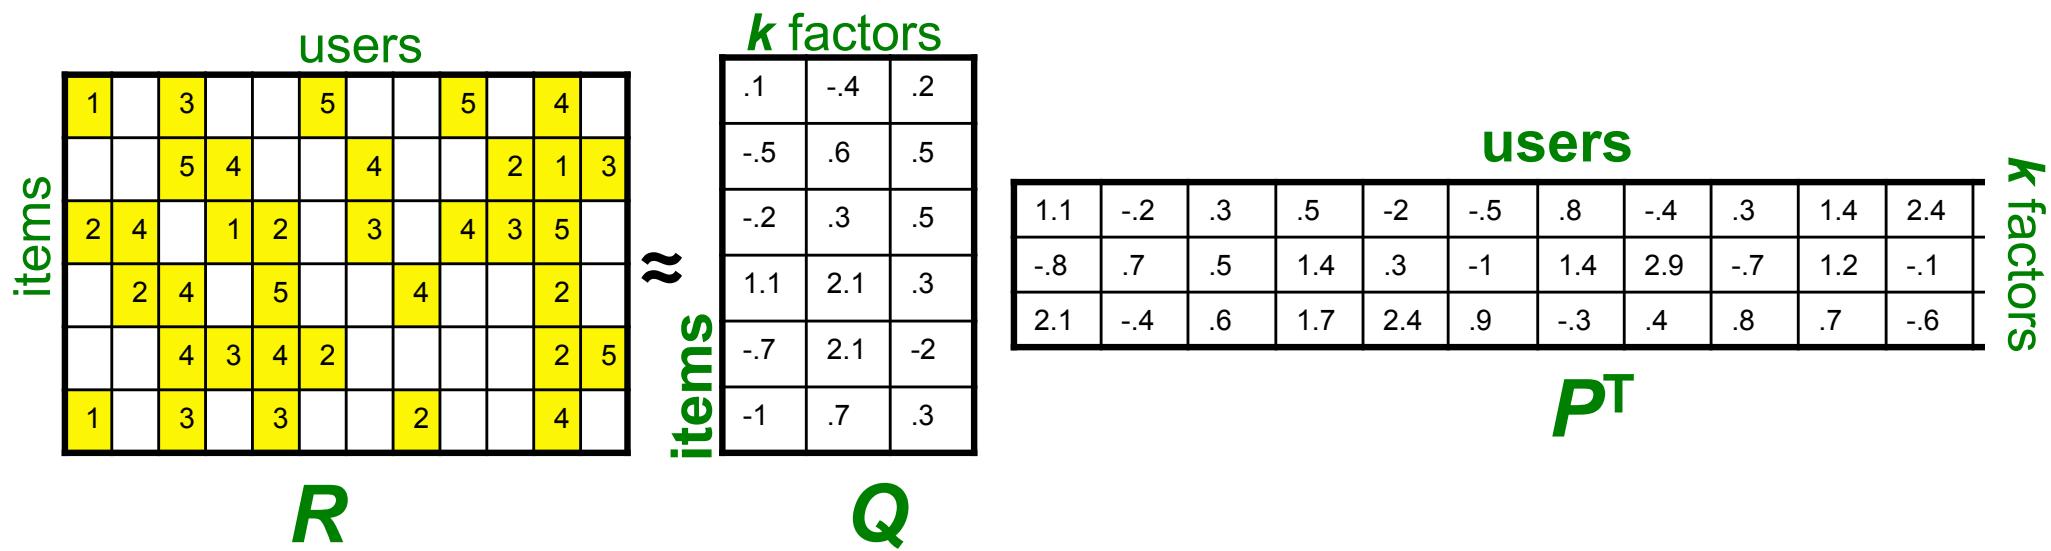
Segment 1

Segment 2

Segment 3

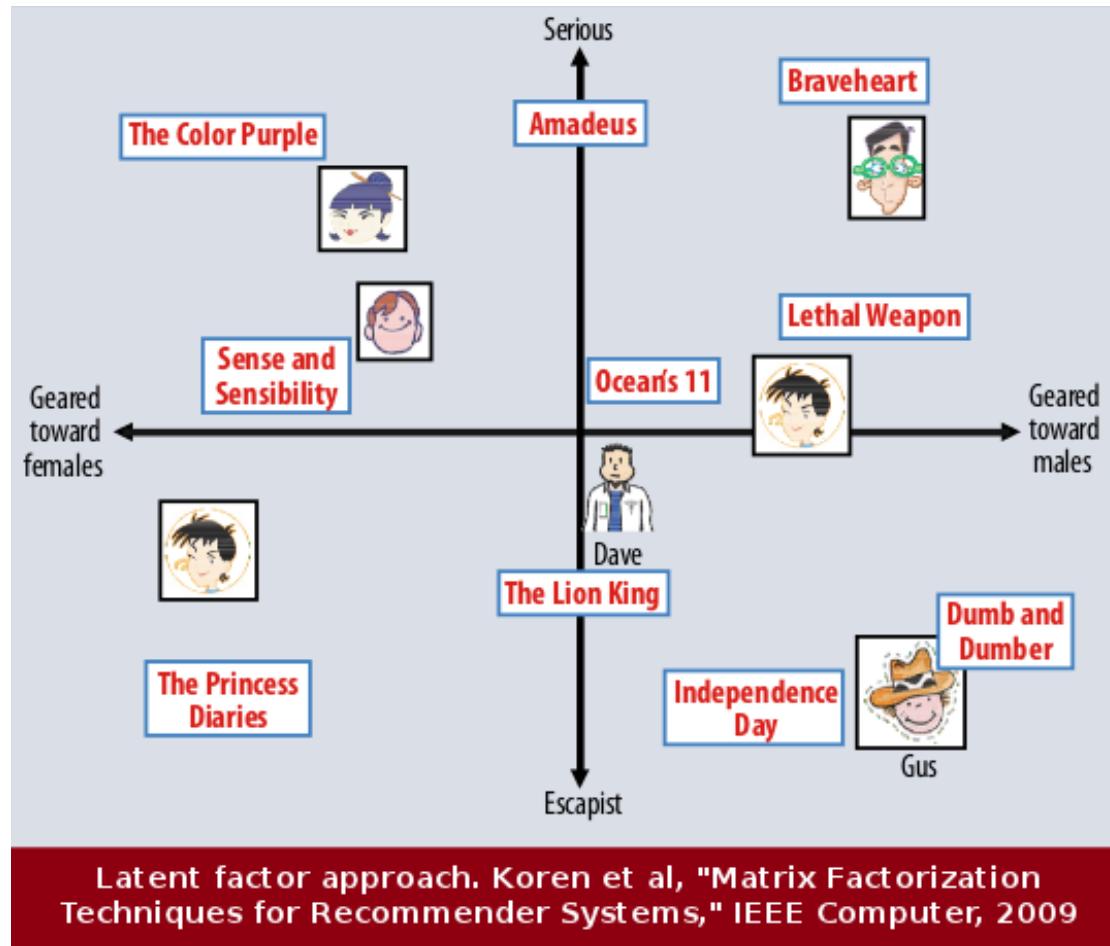
# Latent factors

- IDEA: the utility matrix  $R$  can be factorized, such that  $R = Q \cdot P^T$



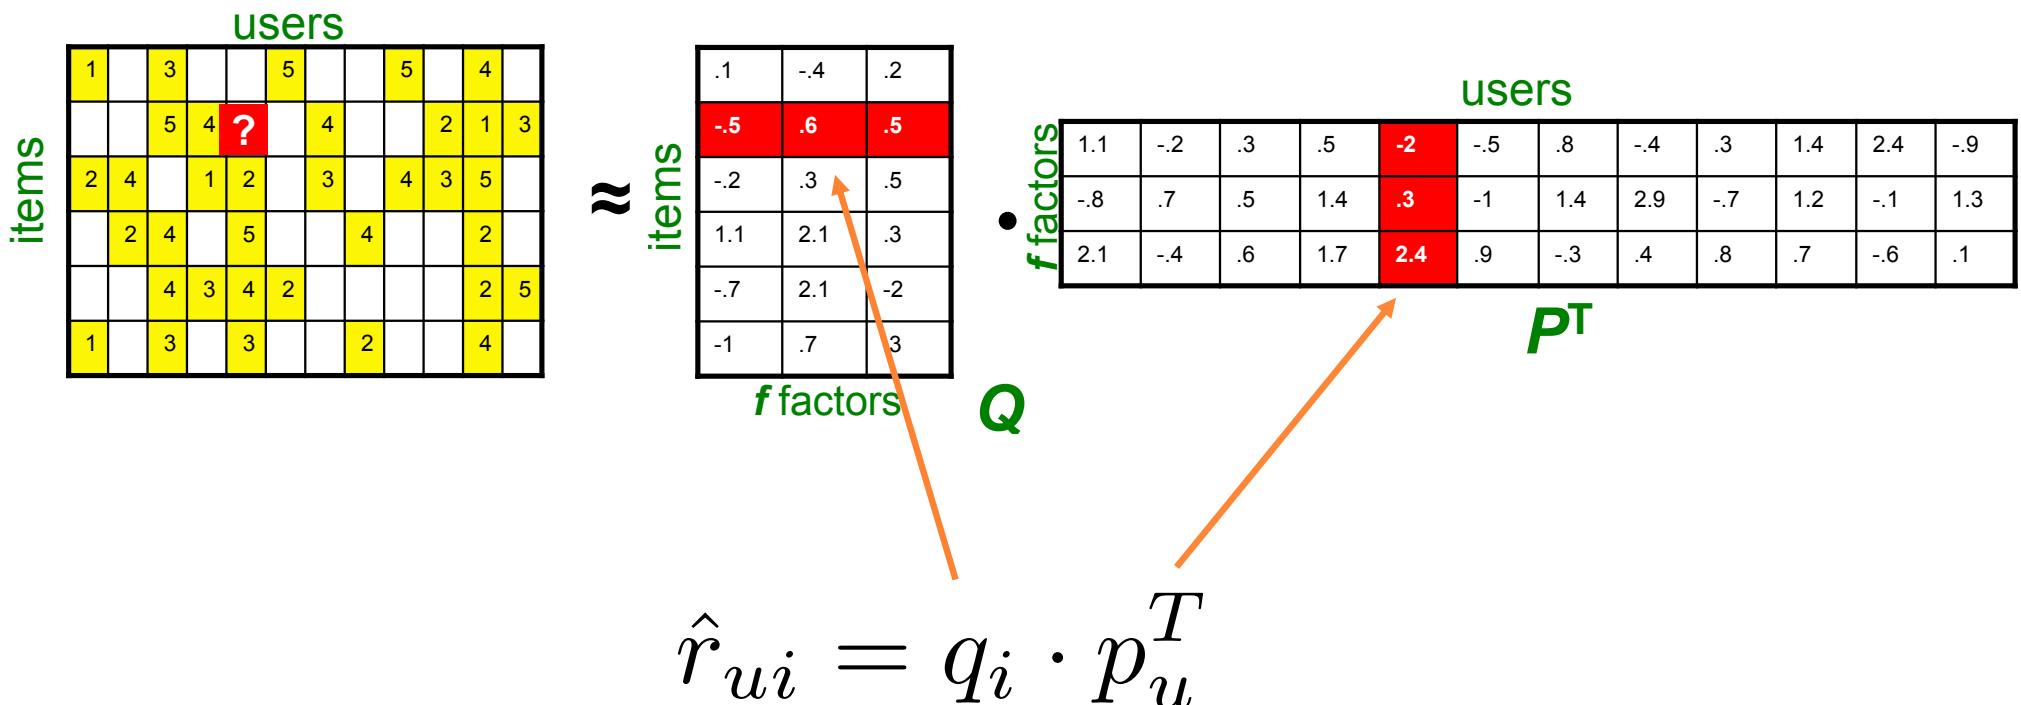
- Factors = new features
- They might not necessarily be intuitive understandable

# Latent factors



# Rating estimation

- How to estimate the missing rating of user  $u$  for item  $i$ ?



# Problem formulation

---

- How to find  $Q$  and  $P^T$ ?
- SVD could be an option  $[U, S, V] = \text{svd}(R)$ 
  - $R = U \cdot S \cdot V^T$
  - $Q = U$  and  $P^T = S \cdot V^T$
  - KO: SVD is not defined when entries are missing
- Regularized solution

$$\min_{q^*, p^*} \sum_{(u,i) \in \kappa} (r_{ui} - q_i^T p_u)^2 + \lambda (\|q_i\|^2 + \|p_u\|^2)$$

which can be minimized using:

- Alternating least squares
- Stochastic gradient descent

# Adding biases

---

- Static global effects

$$\hat{r}_{ui} = \mu + b_i + b_u + q_i^T p_u$$

$$\min_{p^*, q^*, b^*} \sum_{(u,i) \in \kappa} (r_{ui} - \mu - b_u - b_i - p_u^T q_i)^2 + \lambda (\|p_u\|^2 + \|q_i\|^2 + b_u^2 + b_i^2)$$

- Temporal dynamics

$$\hat{r}_{ui}(t) = \mu + b_i(t) + b_u(t) + q_i^T p_u(t)$$

- Rise in the average movie rating (Netflix)
- Users prefer new movies without any reason

# Other approaches

---

## SLIM: Sparse Linear Methods for Top-N Recommender Systems

Xia Ning and George Karypis  
Computer Science & Engineering  
University of Minnesota, Minneapolis, MN  
Email: {xning,karypis@cs.umn.edu}

<http://glaros.dtc.umn.edu/gkhome/node/774>

# Outline

---

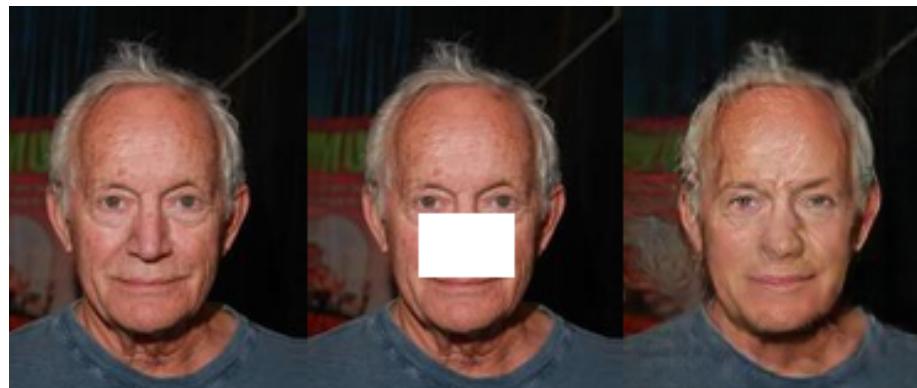
1. Introduction
2. Simple recommenders
3. Content-based recommenders
4. Collaborative Filtering
5. Evaluating recommender systems
6. Latent Factors models
7. **Recommenders using Deep Learning**
8. Recommenders in action (next lecture)



# Deep Learning RS

---

- Restricted Boltzmann Machines for collaborative filtering
  - Generative (autoencoder)
- Autorec
  - Denoising autoencoder



# Deep Learning RS

---

- Deep Matrix Factorization
  - CTR prediction
  - Embeddings

```
#Using keras
emb_u = Embedding(N, K)
emb_i = Embedding(M, K)

u = input()
i = input()

eu = emb_u(u)
ei = emb_i(i)

r = dot(eu, ei)
```

# Deep Learning RS

---

- RNN
  - What comes next?

# Outline

---

1. Introduction
2. Simple recommenders
3. Content-based recommenders
4. Collaborative Filtering
5. Evaluating recommender systems
6. Latent Factors models
7. Recommenders using Deep Learning
- 8. Recommenders in action (next lecture)**



# References

---

- J. Leskovec, A. Rajaraman, J. Ullman. Mining of Massive Datasets
- Recommender Systems, The Textbook. CC Agarwal. Springer 2016.
- Practical Recommender Systems, Kim Falk. Manning 2019.
- Building Recommender Systems with Machine Learning and AI. Frank Kane.
- Recommender Systems and Deep Learning in Python. Lazy Programer Inc., Udemy